

Sumit Kumar

Data Engineer | NIT Patna '21

+91-7549980508

sumit749284@gmail.com

linkedin.com/in/smaxiso

Skills

Languages: Python, Java, Scala

Data: Spark, PySpark, SQL, Kafka, Kafka Connect, Debezium CDC, Airflow

Cloud: AWS (ECS, MSK, DynamoDB, DocumentDB, S3, Bedrock), GCP (BigQuery, Dataproc, Dataflow)

GenAI: OpenAI GPT, AWS Bedrock, LangChain, Prompt Engineering

Frameworks: Spring Boot, FastAPI, React

Tools: Git, Docker, Jenkins, Maven

Experience

Cocoblu

Feb 2026 – Present

Senior Manager – DSIT (Data Science & IT)

Bangalore, India

Data Pipelines & Platform Architecture

- Architected a **React/FastAPI** Master Data Management platform, consolidating **65+ pipelines** from ad-hoc Excel/Glue workflows into a centrally governed system with **RBAC**-scoped self-service access for business teams.
- Re-architected **99 parallel data pipelines** into a unified, config-driven **ETL framework** – cutting infrastructure cost by **90+%** and eliminating the need for custom engineering on new integrations. Built the underlying **SP-API SDK** (SigV4 signing, token caching, circuit breakers), now adopted across the team's automation suite.

Agentic AI Systems

- Architected an AI-powered product classification pipeline across **50,000+ ASINs** using **Gemini 2.5 Flash** and **AWS Bedrock** with automatic model fallback, global image deduplication, and fuzzy-matching – achieving **96.6%** classification accuracy at production scale.
- Architected a **Self-Healing AI Agent** that connects into any underlying automation service, intercepts crashes, generates **LLM-driven** code fixes, validates via **AWS CodeBuild**, and submits Pull Requests for human review.
- Engineered an **agentic pricing pipeline** on **SP-API** that scores product similarity and branches decisioning by product characteristics for continuous, automated competitor-based price actioning. Extended the same architectural approach into **Operations Research** – a hybrid ML/OR system (**ADMM, EOQ**) for supply chain inventory optimization.

Gen Digital (formerly NortonLifeLock)

Dec 2024 – Jan 2026

Data Engineer

Chennai, India

- Engineered custom **Java-based Kafka Connect SMT library** for Debezium CDC processing from AWS DocumentDB, achieving **100% data integrity** and eliminating critical DebeziumException errors across real-time document pipelines.
- Implemented multi-field validation and smart filtering in **Spring Boot** transaction normalization services, **reducing invalid data processing by 40%** and supporting Gen's Unified Data Model initiative on **AWS ECS/MSK**.
- Designed and deployed **Python event-driven microservice** on ECS for member data operations, integrating Kafka for event consumption, DynamoDB for persistence, with comprehensive **GDPR compliance** and GUID sync strategy.
- Architected **Financial Insights Provider** — a Python/FastAPI microservice with **DynamoDB caching (3x speedup)**, 6 insight generators (Fee Monitoring, Duplicate Detection), and intelligent merge-on-PATCH logic for cross-scan deduplication.

FINN

- Led development of a **full-stack module** within the financial wellness platform **leveraging OpenAI GPT** for behavioral analysis, proprietary emotion scoring intelligence, and **Monte Carlo simulations** for retirement planning.

PayPal (via TCS)

Jul 2021 – Nov 2024

Data Engineer

Bangalore, India

Data Migration Framework (Mar 2023 – Nov 2024)

- Developed scalable **ETL framework** for PayPal's data migration using Python, AWS, and BigQuery — achieved **20% faster migration** and **30% improved scalability**.
- Built automated dashboards with **Matplotlib** for stakeholder visibility and deployed orchestration using **Airflow** with auto-generated DAG scripts.

Lynx Entity Linkage Optimization (Jan 2024 – May 2024)

- Optimized **Lynx entity linkage framework** using Scala/Spark and Google's APSS algorithm — achieved **35% better accuracy** and **40% faster** approximate nearest neighbor search via LSH enhancements.

- Conducted comprehensive testing with ML algorithms (RPDBSCAN, K-Means) for similarity scoring validation.
- **On-Demand Merchant Reporting** (*Aug 2021 – Jan 2023*)
- Built **on-demand merchant reporting system** with Python, Oracle DB, and GCP Dataproc — achieved **15% improved accuracy** and **25% faster** report generation.
- Created end-to-end pipeline integrating report requests, authentication, and automated orchestration via internal Airflow (DALM).

Education

National Institute of Technology Patna
Bachelor of Technology in Computer Science & Engineering

2017 – 2021
CGPA: 8.0/10